

Binary Heaps & Priority Queues

Kumkum Saxena

Binary Heaps

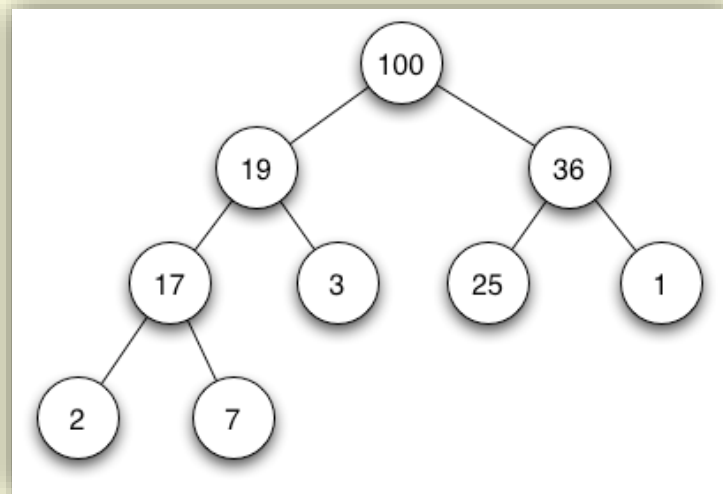
- Heap:

- A heap is an Abstract Data Type
 - Just like stacks and queues are ADTs
 - Meaning, we will define certain behaviors that dictate whether or not a certain data structure is a heap
- So what is a heap?
 - More specifically, what does it do or how do they work?
- A heap looks similar to a tree
 - But a heap has a specific property/invariant that each node in the tree **MUST** follow

Binary Heaps

■ Heap:

- In a heap, all values stored in the subtree of a given node must be less than or equal to the value stored in that node
 - This is known as the heap property




And it is this property that makes a heap a heap!

Binary Heaps

■ Heap:

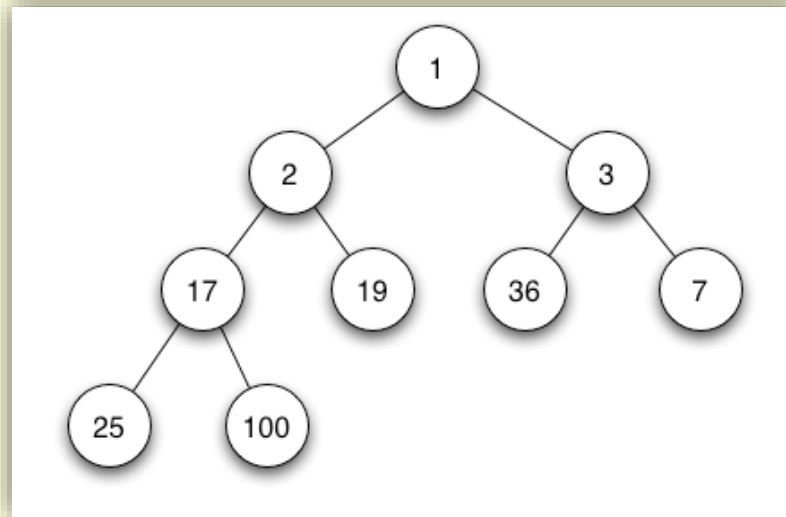
- In a heap, all values stored in the subtree of a given node must be less than or equal to the value stored in that node
 - If B is a child of node A, then the value of node A must be greater than or equal to the value of node B
 - This is called a **Max-Heap**
 - Where the root stores the highest value of any given subtree



Binary Heaps

■ Heap:

- Alternatively, if all values stored in the subtree of a given node are greater than or equal to the value stored in that node
 - This is called a **Min-Heap** (where root is smallest value)




Binary Heaps

- Binary Heap:

- What we just described was a basic Heap
- Now for a heap to be Binary Heap, it must adhere to one other property:
- The **Shape Property**:
 - The heap must be a complete binary tree
 - Meaning, all levels of the tree, except possibly the last one, must be fully filled
 - And if the last level is not complete, the nodes of the level are filled from left to right
 - ***And it just so happens that the previous pictures shown were all examples of binary heaps

Binary Heaps

- Building a Complete Binary Tree:




Root

When a complete binary tree is built, its first node must be the root.

Binary Heaps

- Building a Complete Binary Tree:


Left child
of the
root



The second node is
always the left child
of the root.

Binary Heaps


- Building a Complete Binary Tree:



The third node is
always the right child
of the root.

Binary Heaps


- Building a Complete Binary Tree:



The next nodes always fill the next level from left-to-right.

Binary Heaps


- Building a Complete Binary Tree:



The next nodes always fill the next level from left-to-right.

Binary Heaps


- Building a Complete Binary Tree:



The next nodes always fill the next level from left-to-right.

Binary Heaps


- Building a Complete Binary Tree:



The next nodes always fill the next level from left-to-right.


Binary Heaps

- Building a Complete Binary Tree:



Binary Heaps

- Building a Complete Binary Tree:



This is an example of a **MaxHeap**

Each node in a heap contains a key that can be compared to other nodes' keys.

Binary Heaps

■ Binary Heap:

- New nodes are always added at the lowest level
 - And are inserted from left to right
- There is no particular relationship among the data items in nodes on any given level
 - Even if the nodes have the same parent
 - Example: the right node does not necessarily have to be larger than the left node (as in BSTs)
- The only ordering property for heaps is the one already defined
 - Root of any given subtree is either largest or smallest element in that tree...either a max-heap or a min-heap

Binary Heaps

- Binary Heap:
 - The tree never becomes unbalanced
 - A heap is not a sorted structure
 - But it can be regarded as partially ordered
 - Since the minimum value is always at the root
 - A given set of data can be formed into many different heaps
 - Depending on the order in which the data arrives

Binary Heaps

- Binary Heap:
 - “Okay, great...whupdedoo”
 - Yeah, we now know what a binary heap is
 - But how does it help us?
 - What is its purpose?
- Binary heaps are usually used to implement another abstract data type:
 - A **priority queue**

Binary Heaps

■ Priority Queues:

- A priority queue is basically what it sounds like
 - it is a queue
 - Which means that we will have a line
 - But the first person in line is not necessarily the first person out of line
 - Rather, the **queuing order is based on a priority**
 - Meaning, if one person has a higher priority, that person goes right to the front

■ Examples:

- Emergency room:
 - Higher priority injuries are taken first

Binary Heaps

■ Priority Queues:

■ The model:

- Requests are inserted in the order of arrival
- The request with the highest priority is processed first
 - Meaning, it is removed from the queue
- Priority can be indicated by a number
 - But you have to determine what has most priority
 - Maybe your application results in smallest number having the highest priority
 - Maybe the largest number has the highest priority
 - This really isn't important and is an implementation detail

Binary Heaps

■ Priority Queues:

- So how could we implement a priority queue?
 - Sorted Linked List
 - Higher priority items are ALWAYS at the front of the list
 - Example: a check out line in a supermarket
 - But people who are more important can cut in line
 - Running Time:
 - $O(n)$ insertion time: you have to search through, potentially, n nodes to find the correct spot (based on priority)
 - $O(1)$ deletion time (finding the node with the highest priority) since the highest priority node is first node of the list

Binary Heaps

■ Priority Queues:

- So how could we implement a priority queue?
 - Unsorted Linked List
 - Keep a list of elements as a queue
 - To add an element, append it to the end
 - To remove an element, search through all the elements for the one with the highest priority
 - Running Time:
 - $O(1)$ insertion time: you simply add to the end of the list
 - $O(n)$ deletion time: you have to, potentially, search through all n nodes to find the correct node to delete

Binary Heaps

■ Priority Queues:

- So how could we implement a priority queue?
 - **Correct Method: Binary Heap!**
 - We use a binary heap to implement a priority queue
 - So we are using one abstract data type to implement another abstract data type
 - Running time ends up being $O(\log n)$ for both insertion and deletion into a Heap
 - FindMin (finding the minimum) ends up being $O(1)$
 - cuz we just find (look at) the root, which is $O(1)$
 - So now we look at how to maintain a heap/priority queue
 - How to insert into and delete from a heap
 - And how to build a heap

Binary Heaps

- Adding Nodes to a Binary Heap
 - Assume the existence of a current heap
 - Remember:
 - The binary heap **MUST** follow the Shape property
 - The tree must be balanced
 - Insertions will be made in the **next available spot**
 - Meaning, at the last level
 - and at the next spot, going from left to right
 - But what will most likely happen when you do this?
 - **The Heap property will NOT be maintained**

Binary Heaps

■ Adding Nodes to a Binary Heap

■ Given this Binary Heap:

- And it is a Max-heap

■ We now add a new node


- With data value 42

■ We add at the last position

■ But this voids the Heap Property

- 42 is greater than both 27 and 35

■ So we must fix this!



Binary Heaps

■ Adding Nodes to a Binary Heap


■ **Percolate Up** procedure

- In order to fix the out of place node, we must follow the following “Percolate Up” procedure
 - If the parent of the newly inserted node is less than the newly inserted node (this is clearly for a “max heap”)
 - Then SWAP them
 - This counts as one “Percolate Up” step
 - Continue this process until the new node finds the correct spot
 - Continue SWAPPING until the parent of the new node has a value that is greater than the new node
 - Or if the new node reaches all the way to the root
 - This is now the new “home” for this node

Binary Heaps

■ Adding Nodes to a Binary Heap


- Put the new node in the next available spot.
- Push the new node upward, swapping with its parent until the new node reaches an acceptable location.



Binary Heaps

■ Adding Nodes to a Binary Heap


- Put the new node in the next available spot.
- Push the new node upward, swapping with its parent until the new node reaches an acceptable location.



Binary Heaps

■ Adding Nodes to a Binary Heap


- Put the new node in the next available spot.
- Push the new node upward, swapping with its parent until the new node reaches an acceptable location.



Binary Heaps

■ Adding Nodes to a Binary Heap

- 42 has now reached an acceptable location
- Its parent (node 45) has a value that is greater than 42
- This process is called Percolate Up
- Other books call it Heapification Upward
- What is important is how it works



Binary Heaps

- Adding Nodes to a Binary Heap
 - Percolate Up procedure
 - What is the Big-O running time of insertion into a heap?
 - The actual insertion is simply $O(1)$
 - We simply insert at the last position
 - And you will see (in a bit) how we quick access to this position
 - But when we do this,
 - We need to fix the tree to maintain the Heap Property
 - Percolate Up takes $O(\log n)$ time
 - Why?
 - Because the height of the tree is $\log n$
 - Worst case scenario is having to SWAP all the way to the root
 - **So the overall running time of an insertion is $O(\log n)$**

Binary Heaps

- Deleting Nodes from a Binary Heap
 - We will write a function called deleteMin (or deleteMax)
 - Which node will we ALWAYS be deleting?
 - Remember:
 - We are using a Heap to implement a priority queue!
 - And in a priority queue, we always delete the first element
 - The one with the highest priority
 - So we will ALWAYS be deleting the ROOT of the tree
 - So this is quite easy!
 - deleteMin (or deleteMax for a Max Heap) simply deletes the root and returns its value to main

Binary Heaps

- Deleting Nodes from a Binary Heap
 - We will write a function called deleteMin
 - deleteMin simply deletes the root and returns its value to main
 - But what will happen when we delete the root?
 - We will have a tree with no root!
 - The root will be missing
 - So clearly this needs to be fixed

Binary Heaps

This process is
for a Max-heap

■ Deleting Nodes from a Binary Heap


■ Fixing the tree after deleting the root:

- 1) Copy the last node of the tree into the position of the root
 - 2) Then remove that last node (to avoid duplicates)
 - Note: **The new root is almost assuredly out of place**
 - Most likely, one, or both, of its children will have a greater value than it
 - If so:
 - 3) Swap the new root node with the **greater** of its child nodes
 - This is considered one “**Percolate Down**” step
- Continue this process until the “last node” ends up in a spot where its children have values smaller than it
- Neither child can have a value greater than it

Binary Heaps

■ Deleting Nodes from a Binary Heap


- Given the following Heap:
- We perform a delete
- Which means 45 will get deleted



Binary Heaps


■ Deleting Nodes from a Binary Heap

- Given the following Heap:
- We perform a delete
- Which means 45 will get deleted



Binary Heaps


- Deleting Nodes from a Binary Heap
- The last node now gets moved to the root
- So 27 goes to the root



Binary Heaps

■ Deleting Nodes from a Binary Heap

- The last node now gets moved to the root
- So 27 goes to the root
- 27 is now out of place
- We must Percolate Down




Binary Heaps

■ Deleting Nodes from a Binary Heap

■ Percolate Down:

- Push the out-of-place node downward,
 - swapping with its **larger** child
- until the out-of-place node reaches an acceptable location




Binary Heaps

- Deleting Nodes from a Binary Heap

- **Percolate Down:**

- Push the out-of-place node downward,
 - swapping with its **larger** child
- until the out-of-place node reaches an acceptable location




Binary Heaps

- Deleting Nodes from a Binary Heap

- **Percolate Down:**

- Push the out-of-place node downward,
 - swapping with its **larger** child
- until the out-of-place node reaches an acceptable location




Binary Heaps

■ Deleting Nodes from a Binary Heap

■ Percolate Down:

- 27 has reached an acceptable location
- Its lone child (19) has a value that is less than 27
- So we stop the Percolate Down procedure at this point



Binary Heaps

■ Deleting Nodes from a Binary Heap

- What is the Big-O running time of deletion from a heap?
- The actual deletion itself is $O(1)$
 - cause the minimum value is at the root
 - and we can delete the root of a tree in $O(1)$ time
- But now we need to fix the tree
 - Moving the last node to the root is an $O(1)$ step
 - But then we need to Percolate Down
- Percolate Down takes $O(\log n)$
 - Why?
 - Because the height of the tree is $\log n$
 - And the worst case scenario is having to SWAP all the way to the farthest leaf
- **So the overall running time of a deletion is $O(\log n)$**