


# Red Black Trees

*Kumkum Saxena*

# Red-Black Tree

---

- Red-Black Tree is a self-balancing Binary Search Tree (BST) where every node follows following rules.
  - **1)** Every node has a color either red or black.
  - **2)** Root of tree is always black.
  - **3)** There are no two adjacent red nodes (A red node cannot have a red parent or red child).
  - **4)** Every path from a node (including root) to any of its descendant NULL node has the same number of black nodes.



# Why Red-Black Trees?

---

- Most of the BST operations (e.g., search, max, min, insert, delete.. etc) take  $O(h)$  time where  $h$  is the height of the BST.
- The cost of these operations may become  $O(n)$  for a skewed Binary tree.
- If we make sure that height of the tree remains  $O(\log n)$  after every insertion and deletion, then we can guarantee an upper bound of  $O(\log n)$  for all these operations. The height of a Red-Black tree is always  $O(\log n)$  where  $n$  is the number of nodes in the tree.

# Comparison with AVL Tree

---

- The AVL trees are more balanced compared to Red-Black Trees, but they may cause more rotations during insertion and deletion.
- So if your application involves many frequent insertions and deletions, then Red Black trees should be preferred.
- And if the insertions and deletions are less frequent and search is a more frequent operation, then AVL tree should be preferred over Red-Black Tree.


# How does a Red-Black Tree ensure balance?

---


- A simple example to understand balancing is, a chain of 3 nodes is not possible in the Red-Black tree. We can try any combination of colours and see all of them violate Red-Black tree property.

A chain of 3 nodes is not possible in Red-Black Trees.

Following are **NOT** Red-Black Trees



Following are different possible Red-Black Trees with above 3 keys



# *Black Height of a Red-Black Tree*

---

- *Black height is number of black nodes on a path from root to a leaf.*
- *Leaf nodes are also counted black nodes.*
- *From above properties 3 and 4, we can derive, a **Red-Black Tree of height  $h$  has black-height  $\geq h/2$ .***
- ***Every Red Black Tree with  $n$  nodes has height  $\leq 2\log_2(n+1)$***



- This can be proved using following facts:
  - For a general Binary Tree, let  $k$  be the minimum number of nodes on all root to NULL paths, then  $n \geq 2^k - 1$  (Ex. If  $k$  is 3, then  $n$  is atleast 7). This expression can also be written as  $k \leq \log_2(n+1)$
  - From property 4 of Red-Black trees and above claim, we can say in a Red-Black Tree with  $n$  nodes, there is a root to leaf path with at-most  $\log_2(n+1)$  black nodes.
  - From property 3 of Red-Black trees, we can claim that the number black nodes in a Red-Black tree is at least  $\lfloor n/2 \rfloor$  where  $n$  is the total number of nodes.
- From above 2 points, we can conclude the fact that Red Black Tree with  $n$  nodes has height  $\leq 2\log_2(n+1)$

# Red-Black Tree(Insert)

---

- In AVL tree insertion, we used rotation as a tool to do balancing after insertion caused imbalance. In Red-Black tree, we use two tools to do balancing.

- Recoloring

- Rotation

We try recoloring first, if recoloring doesn't work, then we go for rotation

- 
- The algorithm has mainly two cases depending upon the color of uncle.
  - If uncle is red, we do recoloring. If uncle is black, we do rotations and/or recoloring.
  - Color of a NULL node is considered as BLACK.

Let x be the newly inserted node.

1. Perform standard BST insertion and make the color of newly inserted nodes as RED.
2. If x is root, change color of x as BLACK (Black height of complete tree increases by 1).
3. Do following if color of x's parent is not BLACK **and** x is not root.
  1. **a) If x's uncle is RED** (Grand parent must have been black from property 4)
  2. **.b) If x's uncle is BLACK,**

---


**a) If x's uncle is RED** (Grand parent must have been black from property 4)

**(i)** Change color of parent and uncle as BLACK.

**(ii)** color of grand parent as RED.

**(iii)** Change  $x = x$ 's grandparent, repeat steps 2 and 3 for new  $x$ .

## Uncle Red




x: Current Node, p: Parent, u: Uncle, g: Grandparent

T1, T2, T3, T4 and T5 are subtrees

- **b) If x's uncle is BLACK**, then there can be four configurations for x, x's parent (**p**) and x's grandparent (**g**) (This is similar to AVL Tree)
  - i) Left Left Case** (p is left child of g and x is left child of p)
  - ii) Left Right Case** (p is left child of g and x is right child of p)
  - iii) Right Right Case** (Mirror of case i)
  - iv) Right Left Case** (Mirror of case ii)

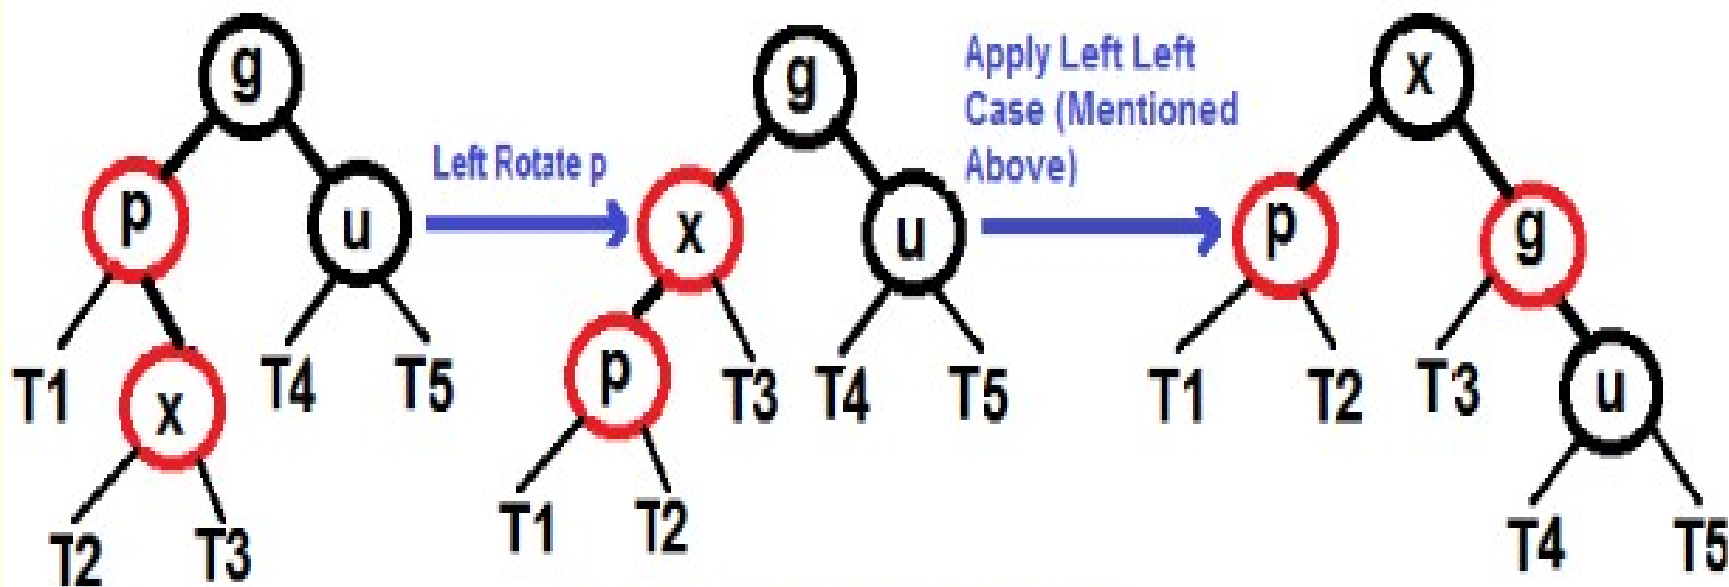
# Left Left Case (See g, p and x)





# Left Right Case (See g, p and x)


## Uncle Black and Left Right Case



x: Current Node, p: Parent, u: Uncle, g: G

T1, T2, T3, T4 and T5 are subtrees

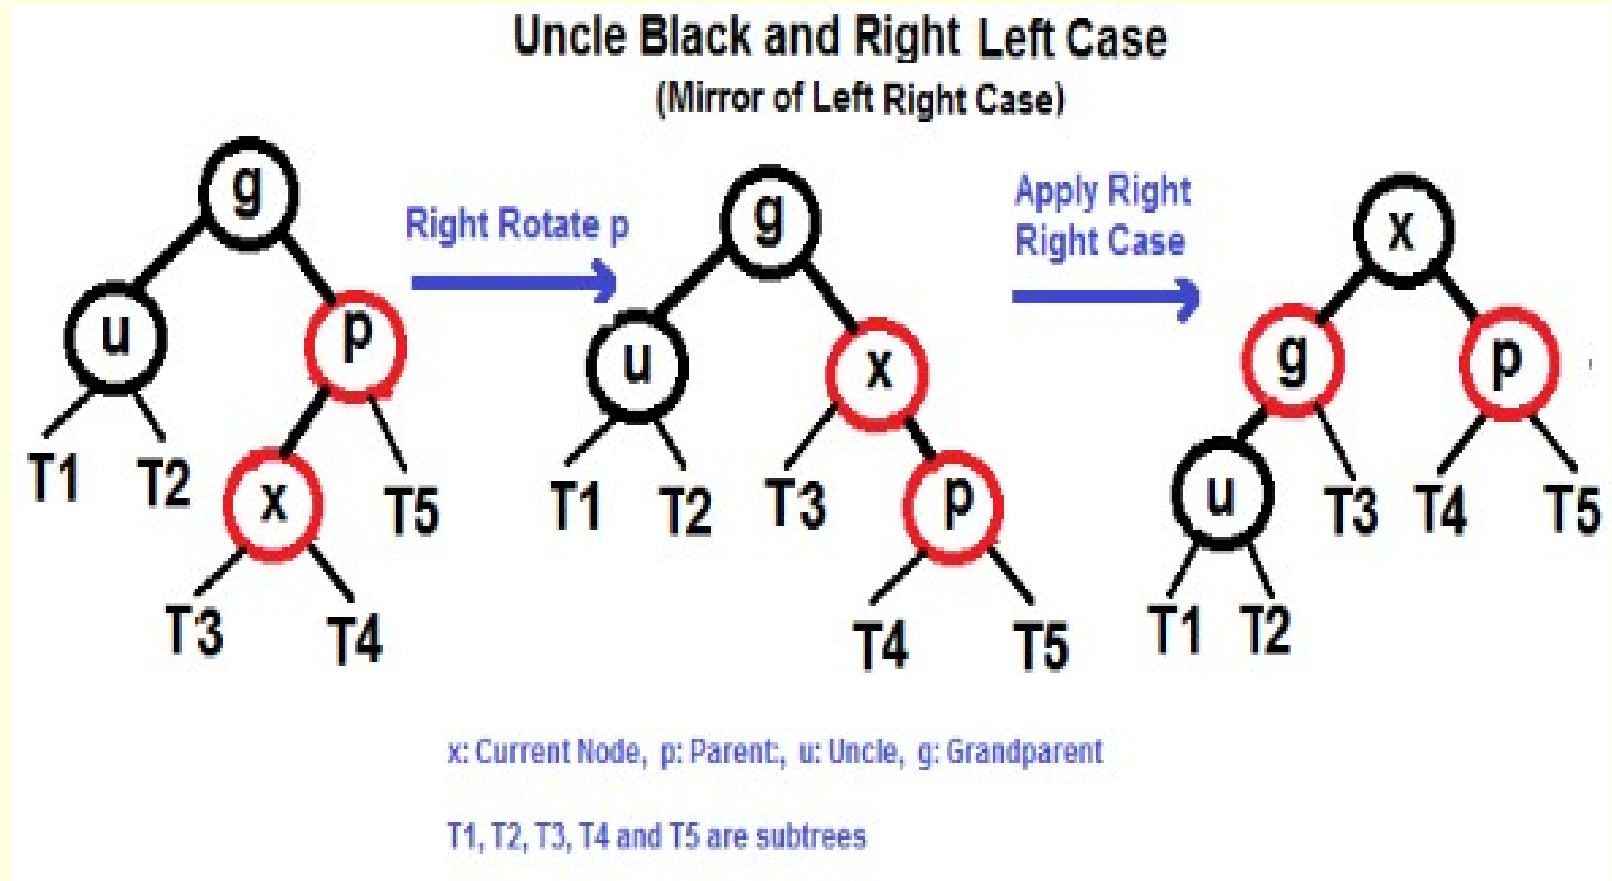
# Right Right Case (See g, p and x)



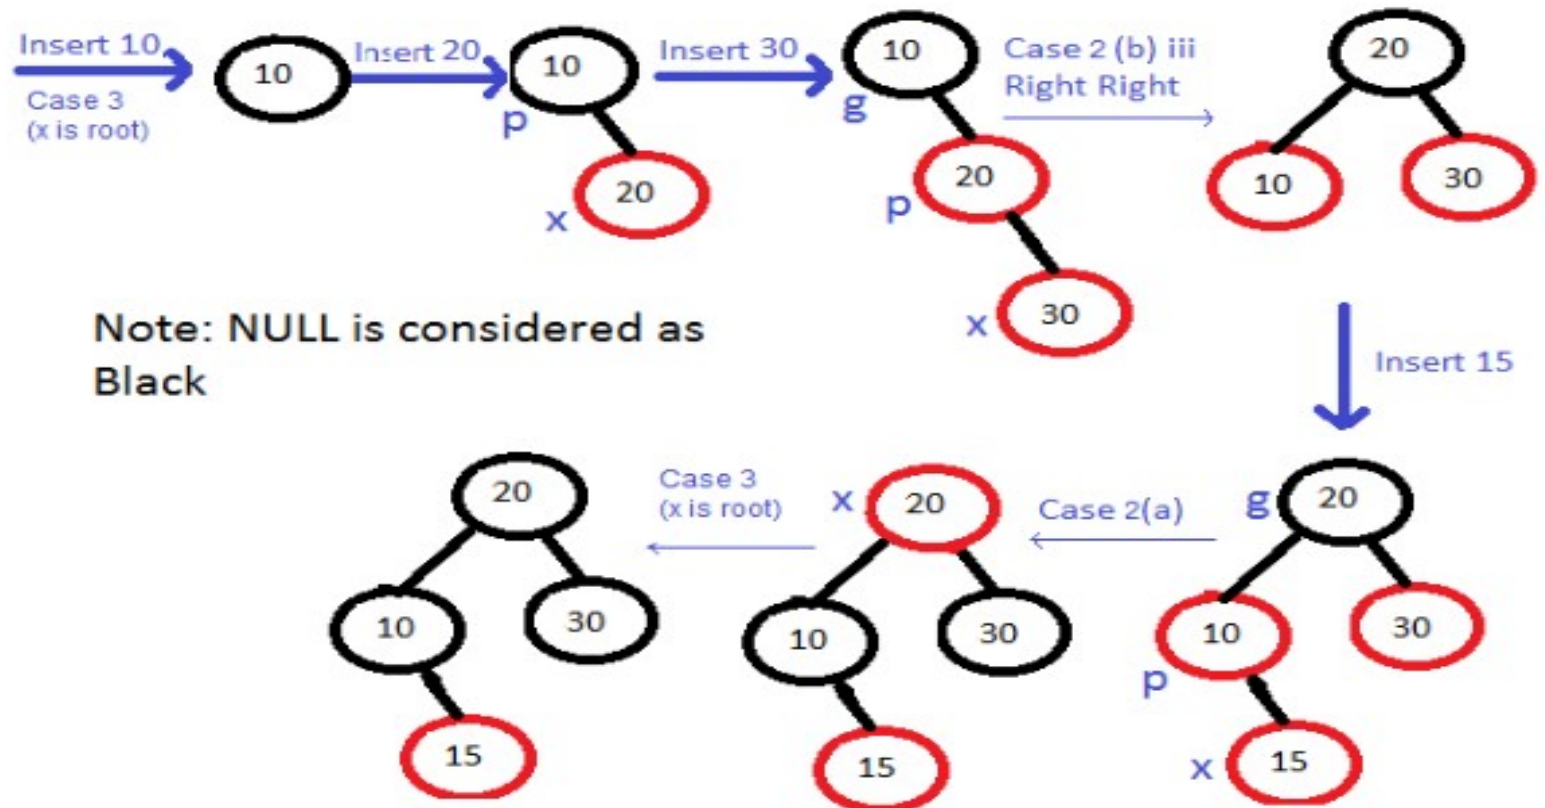
x: Current Node, p: Parent, u: Uncle, g: Grandparent

T1, T2, T3, T4 and T5 are subtrees

# Right Left Case (See g, p and x)



## Insert 10, 20, 30 and 15 in an empty tree



# Algorithm Analysis

---

- $O(\lg n)$  time to get through RB-Insert up to the call of RB-Insert-Fixup.
- Within **RB-Insert-Fixup**:
  - Each iteration takes  $O(1)$  time.
  - Each iteration but the last **moves  $z$  up 2 levels**.
  - $O(\lg n)$  levels  $\Rightarrow O(\lg n)$  time.
  - Thus, insertion in a red-black tree takes  $O(\lg n)$  time.
  - Note: there are at most 2 rotations overall.


# RB Properties Affected by Insert

1. Every node is either **red** or **black** OK!
2. The root is **black** If z is the root  
⇒ **not OK**
3. Every leaf (NIL) is **black** OK!
4. If a node is red, then both its children are black

If  $p(z)$  is red ⇒ **not OK**  
z and  $p(z)$  are both

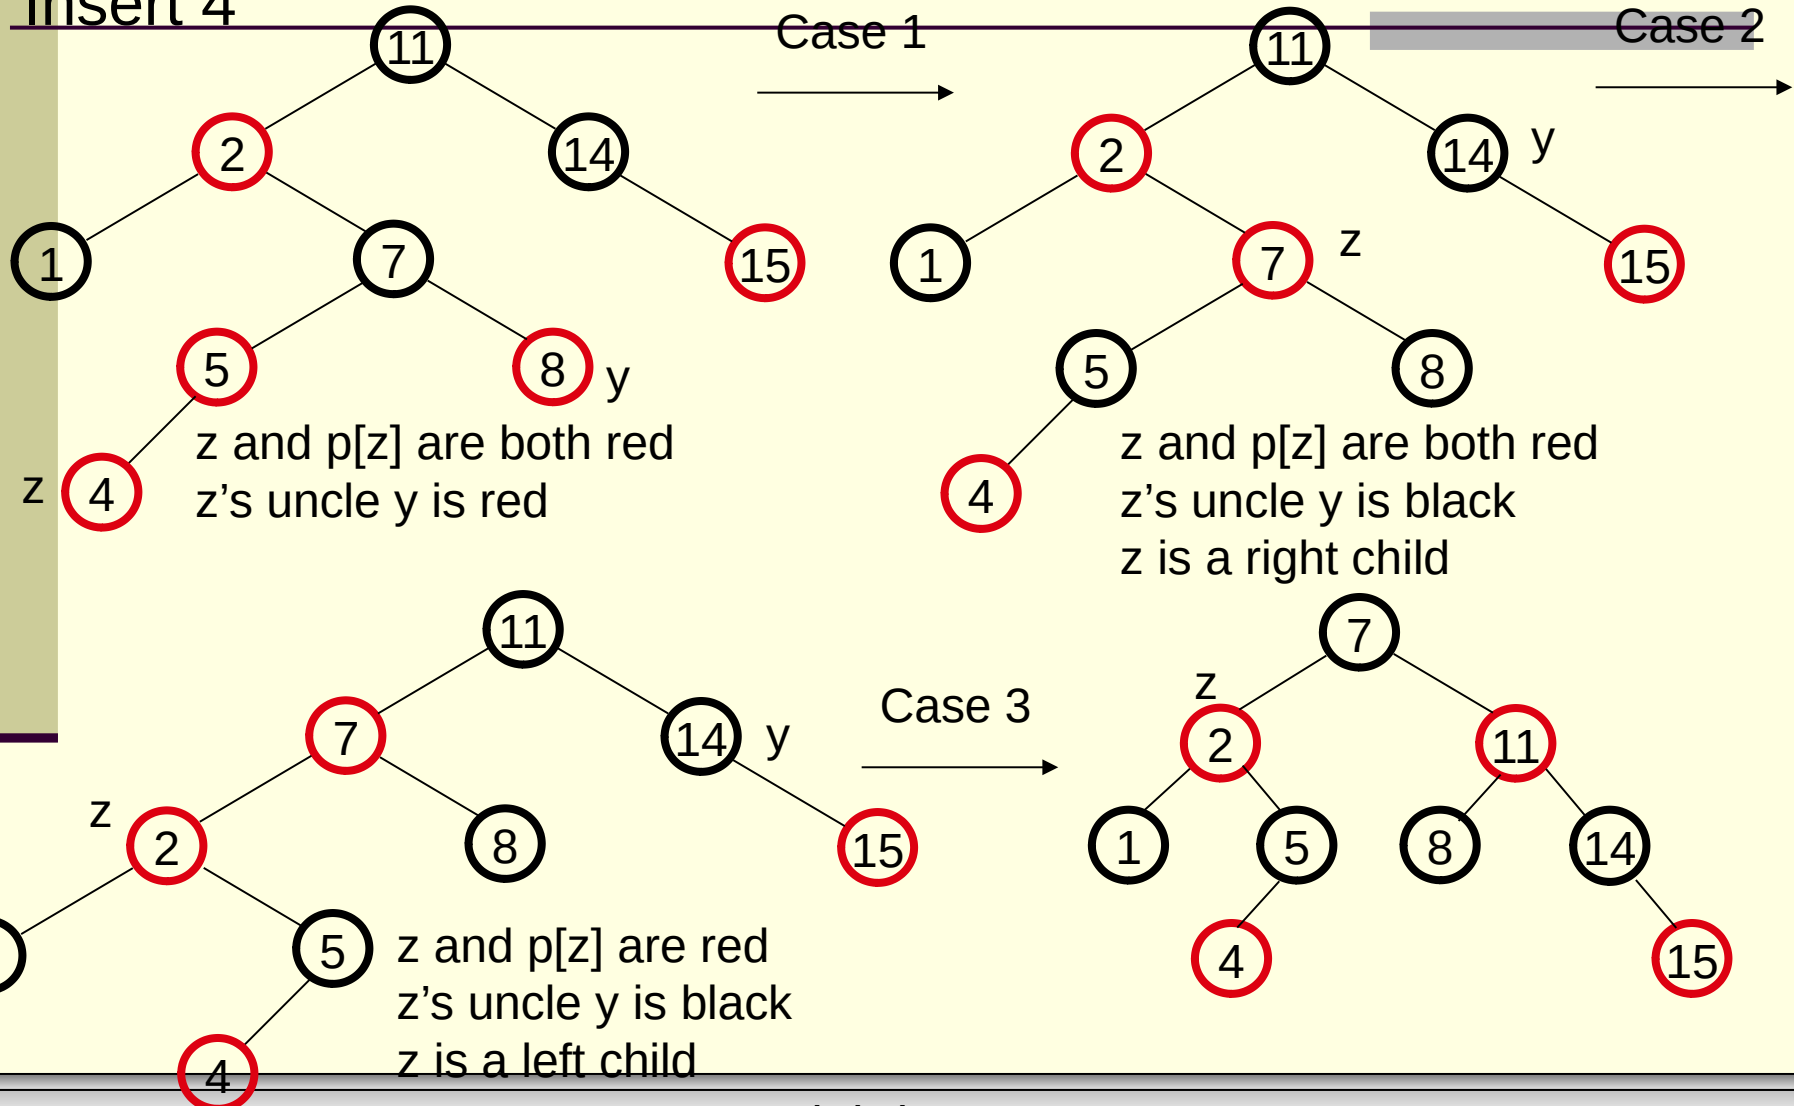
red

- OK!
5. For each node, all paths  
from the node to descendant  
leaves contain the same number  
of black nodes




# Example

Insert 4




# Example of Inserting Sorted Numbers



Insert 1. A leaf so red. Realize it is root so recolor to black.




# Insert 2



make 2 red. Parent  
is black so done.

# Insert 3

Insert 3. Parent is red.  
Parent's sibling is black  
(null) 3 is outside relative  
to grandparent. Rotate  
parent and grandparent




# Insert 4

On way down see  
2 with 2 red  
children.

Recolor 2 red and  
children black.


Realize 2 is root  
so color back to  
black

When adding 4  
parent is black  
so done.




# Insert 5

5's parent is red.  
Parent's sibling is  
black (null). 5 is  
outside relative to  
grandparent (3) so rotate  
parent and grandparent then  
recolor




# Finish insert of 5




# Insert 6

On way down see  
4 with 2 red  
children. Make  
4 red and children  
black. 4's parent is  
black so no problem.




# Finishing insert of 6

6's parent is black  
so done.




# Insert 7

7's parent is red.  
Parent's sibling is  
black (null). 7 is  
outside relative to  
grandparent (5) so  
rotate parent and  
grandparent then recolor






# Finish insert of 7



# Insert 8


On way down see 6  
with 2 red children.

Make 6 red and  
children black. This  
creates a problem  
because 6's parent, 4, is  
also red. Must perform  
rotation.




# Still Inserting 8

Recolored now  
need to  
rotate

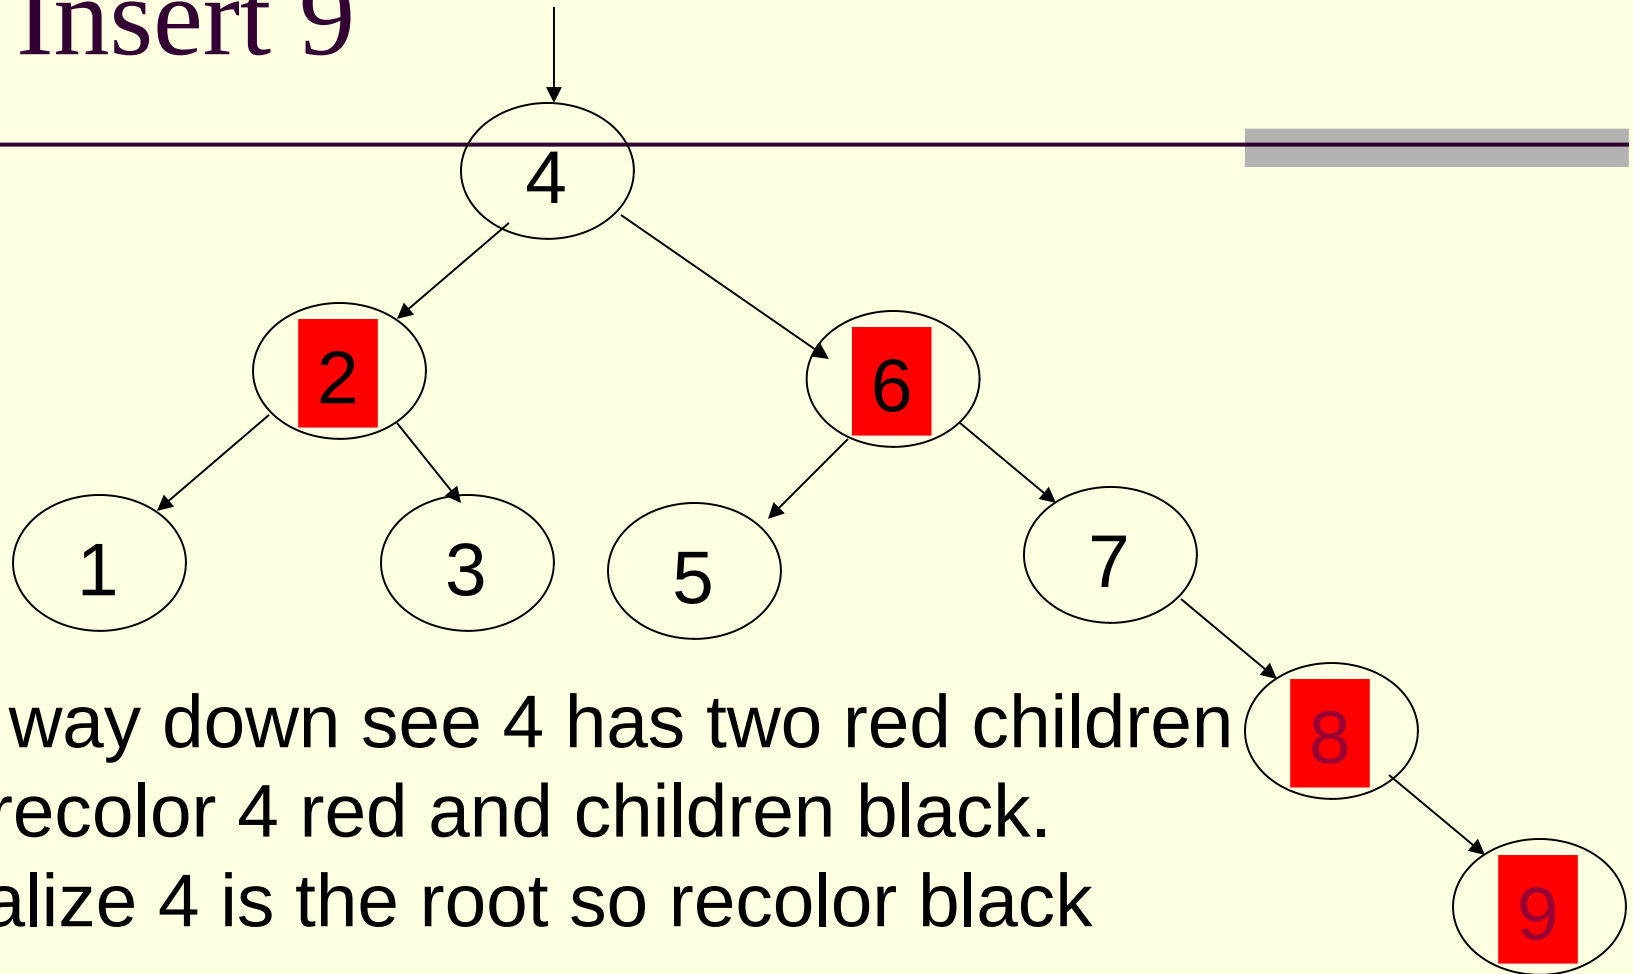


# Finish inserting 8

Recolored  
now  
need to  
rotate




# Insert 9



On way down see 4 has two red children  
so recolor 4 red and children black.


Realize 4 is the root so recolor black

# Finish Inserting 9




After rotations and recoloring

# Insert 10




# Insert 11





# Finish inserting 11



# Red-Black Tree (Delete)

---

- Like Insertion, recoloring and rotations are used to maintain the Red-Black properties.
- In insert operation, we check color of uncle to decide the appropriate case. In delete operation, ***we check color of sibling*** to decide the appropriate case.

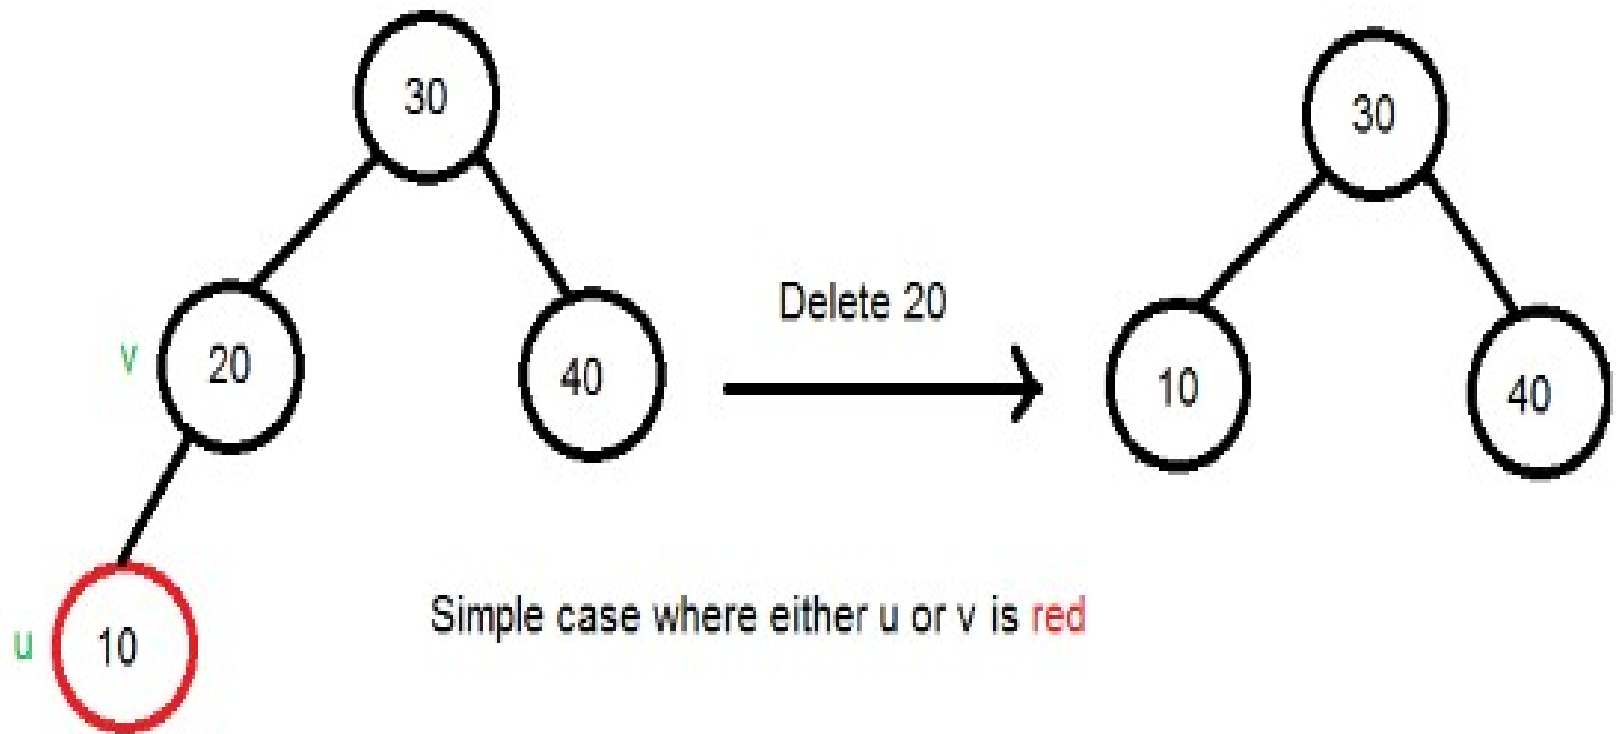
- The main property that violates after insertion is two consecutive reds. In delete, the main violated property is, change of black height in subtrees as deletion of a black node may cause reduced black height in one root to leaf path.
- Deletion is fairly complex process. To understand deletion, notion of double black is used. When a black node is deleted and replaced by a black child, the child is marked as **double black**. The main task now becomes to convert this double black to single black.

# Deletion Steps

---

- 1) Perform standard BST delete.
- When we perform standard delete operation in BST, we always end up deleting a node which is either leaf or has only one child (For an internal node, we copy the successor and then recursively call delete for successor, successor is always a leaf node or a node with one child). So we only need to handle cases where a node is leaf or has one child. Let  $v$  be the node to be deleted and  $u$  be the child that replaces  $v$  (Note that  $u$  is NULL when  $v$  is a leaf and color of NULL is considered as Black).

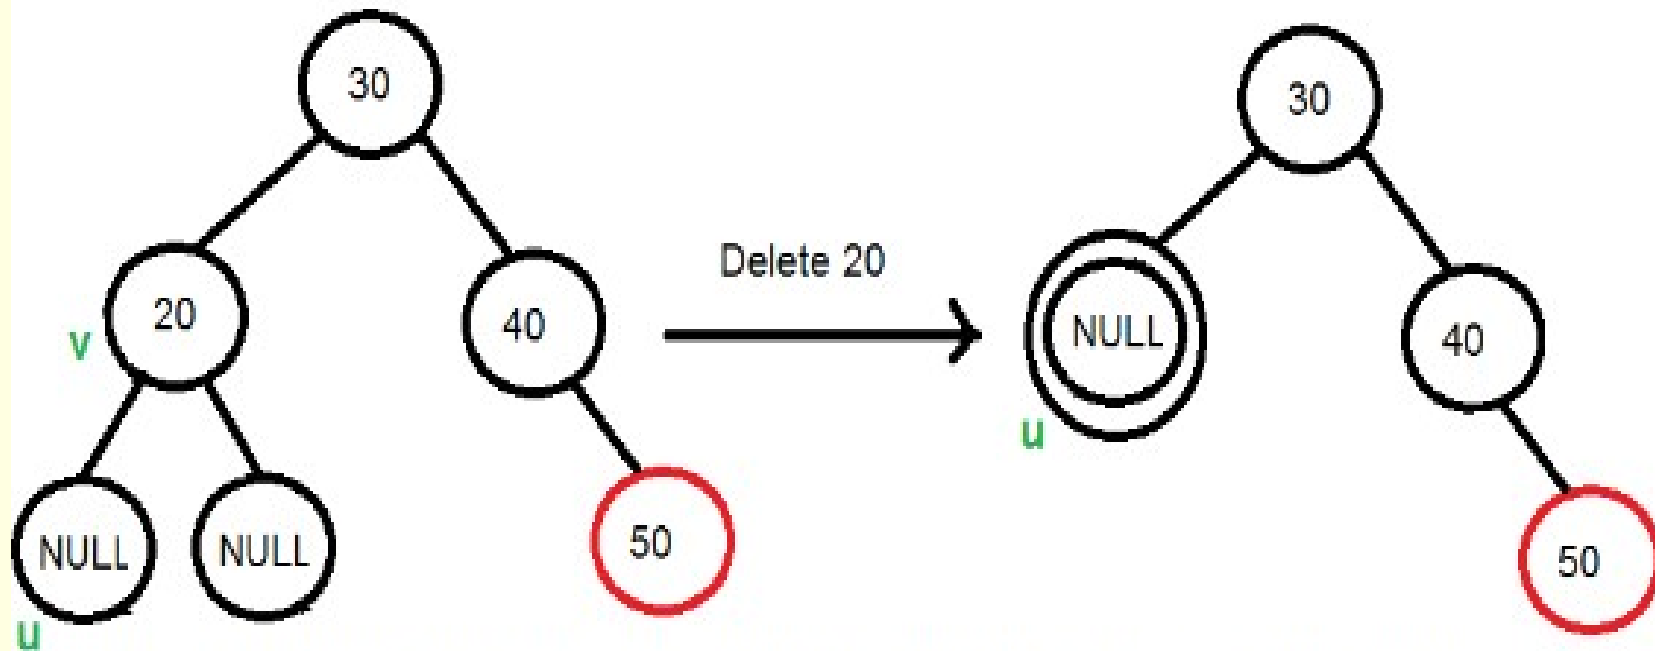
- 
- **2) Simple Case:** If either  $u$  or  $v$  is red, we mark the replaced child as black (No change in black height). Note that both  $u$  and  $v$  cannot be red as  $v$  is parent of  $u$  and two consecutive reds are not allowed in red-black tree.



### 3) If Both u and v are Black.

---

- **3.1)** Color u as double black. Now our task reduces to convert this double black to single black. Note that If v is leaf, then u is NULL and color of NULL is considered as black. So the deletion of a black leaf also causes a double black.

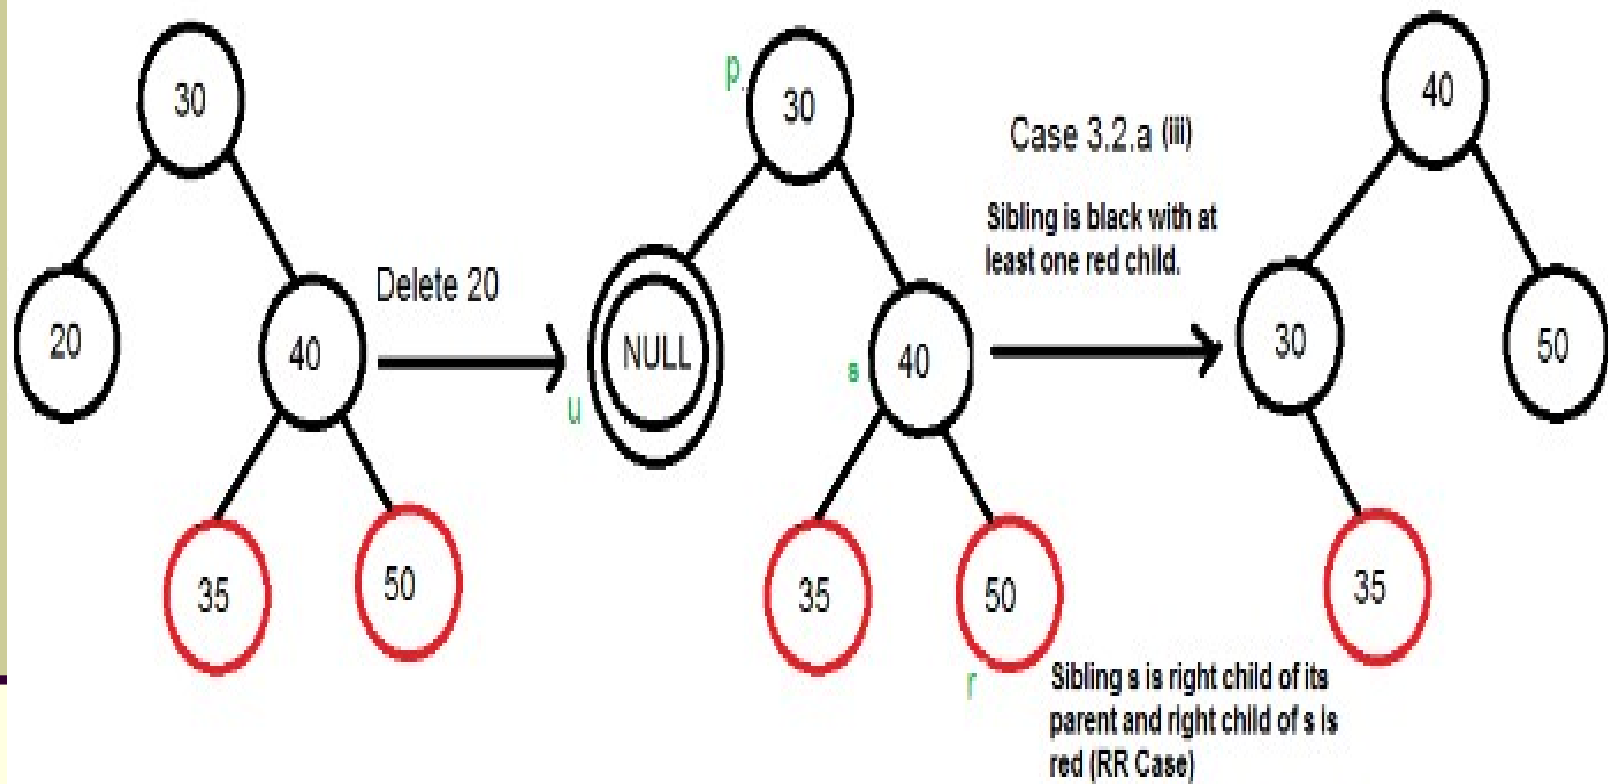


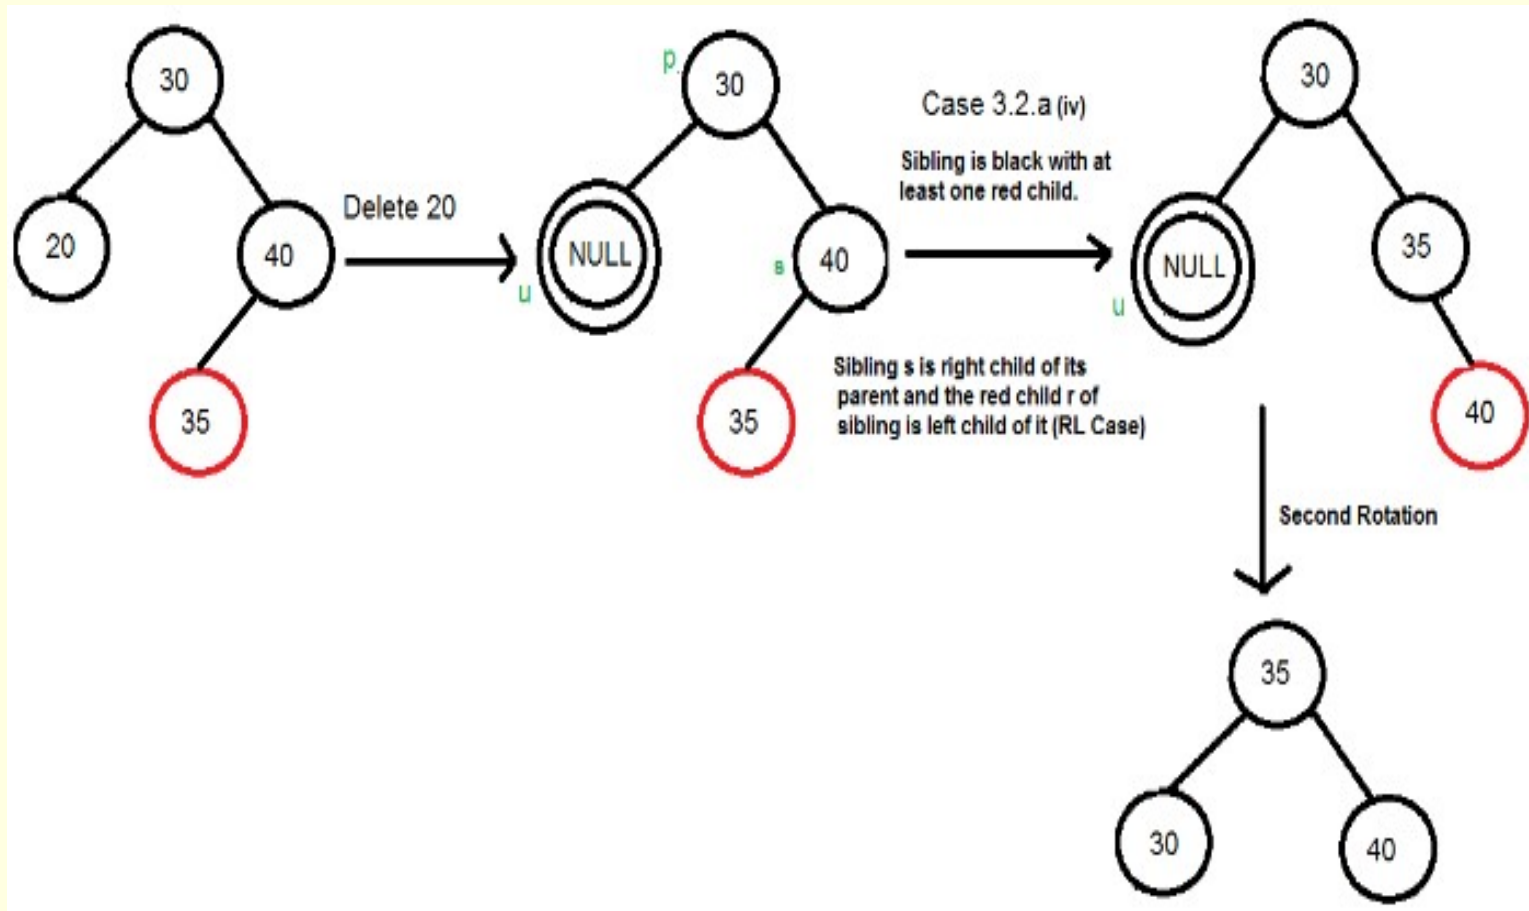
When 20 is deleted, it is replaced by a NULL, so the NULL becomes double black.  
Note that deletion is not done yet, this double black must become single black



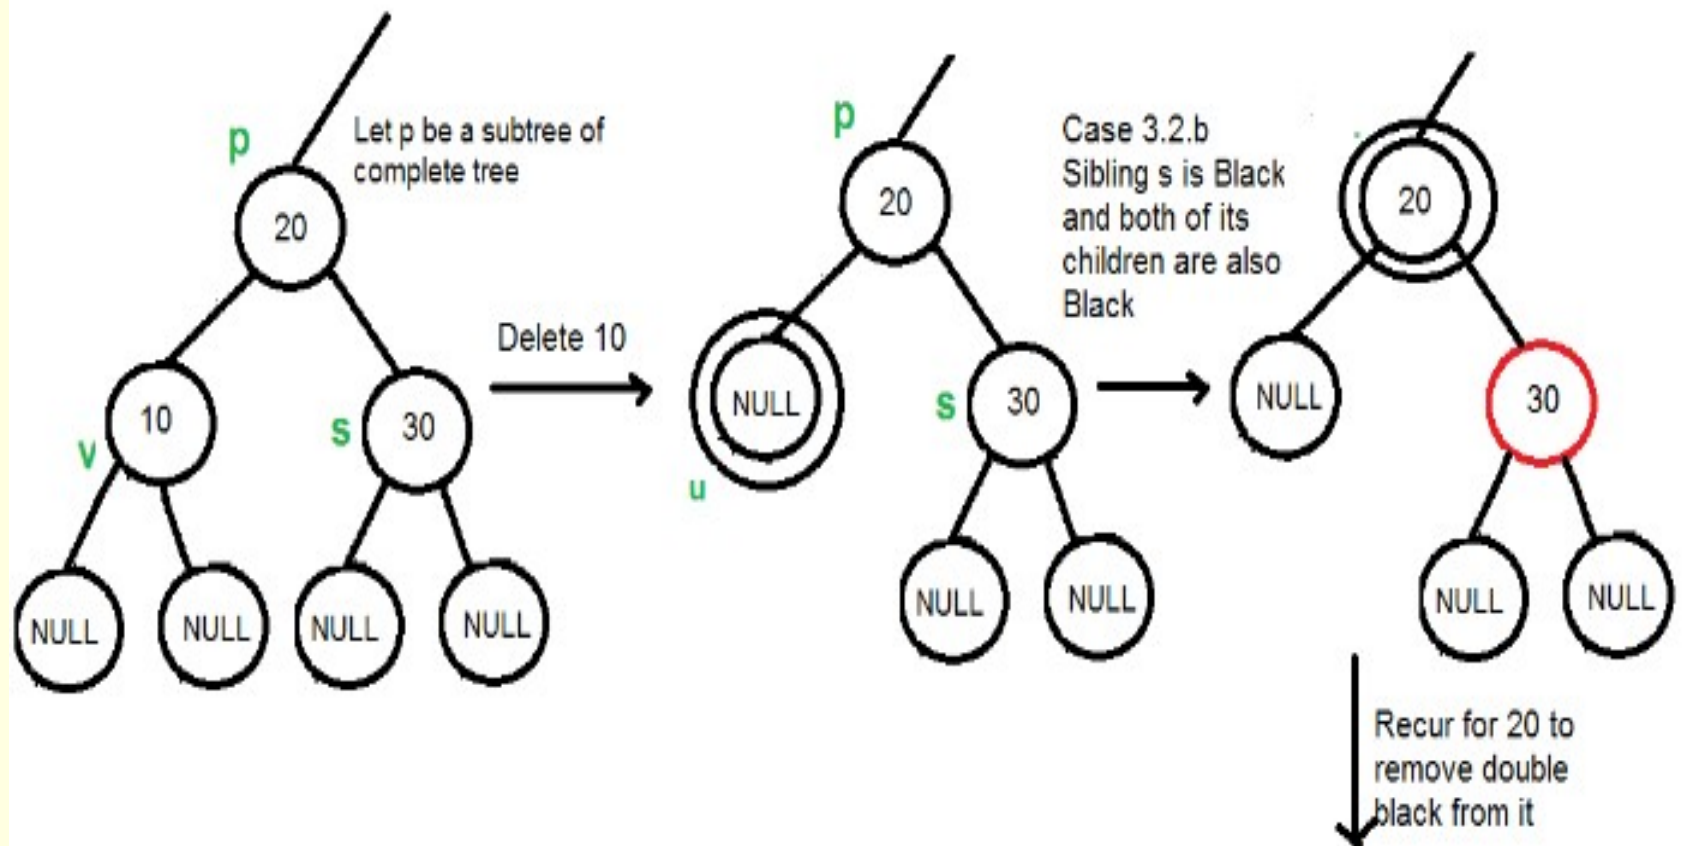
- 
- **3.2)** Do following while the current node  $u$  is double black and it is not root. Let sibling of node be  $s$ .

- **(a): If sibling  $s$  is black and at least one of sibling's children is red**, perform rotation(s). Let the red child of  $s$  be  $r$ . This case can be divided in four subcases depending upon positions of  $s$  and  $r$ .
- **(i) Left Left Case** ( $s$  is left child of its parent and  $r$  is left child of  $s$  or both children of  $s$  are red).
- **(ii) Left Right Case** ( $s$  is left child of its parent and  $r$  is right child).
- **(iii) Right Right Case** ( $s$  is right child of its parent and  $r$  is right child of  $s$  or both children of  $s$  are red)
- **(iv) Right Left Case** ( $s$  is right child of its parent and  $r$  is left child of  $s$ )

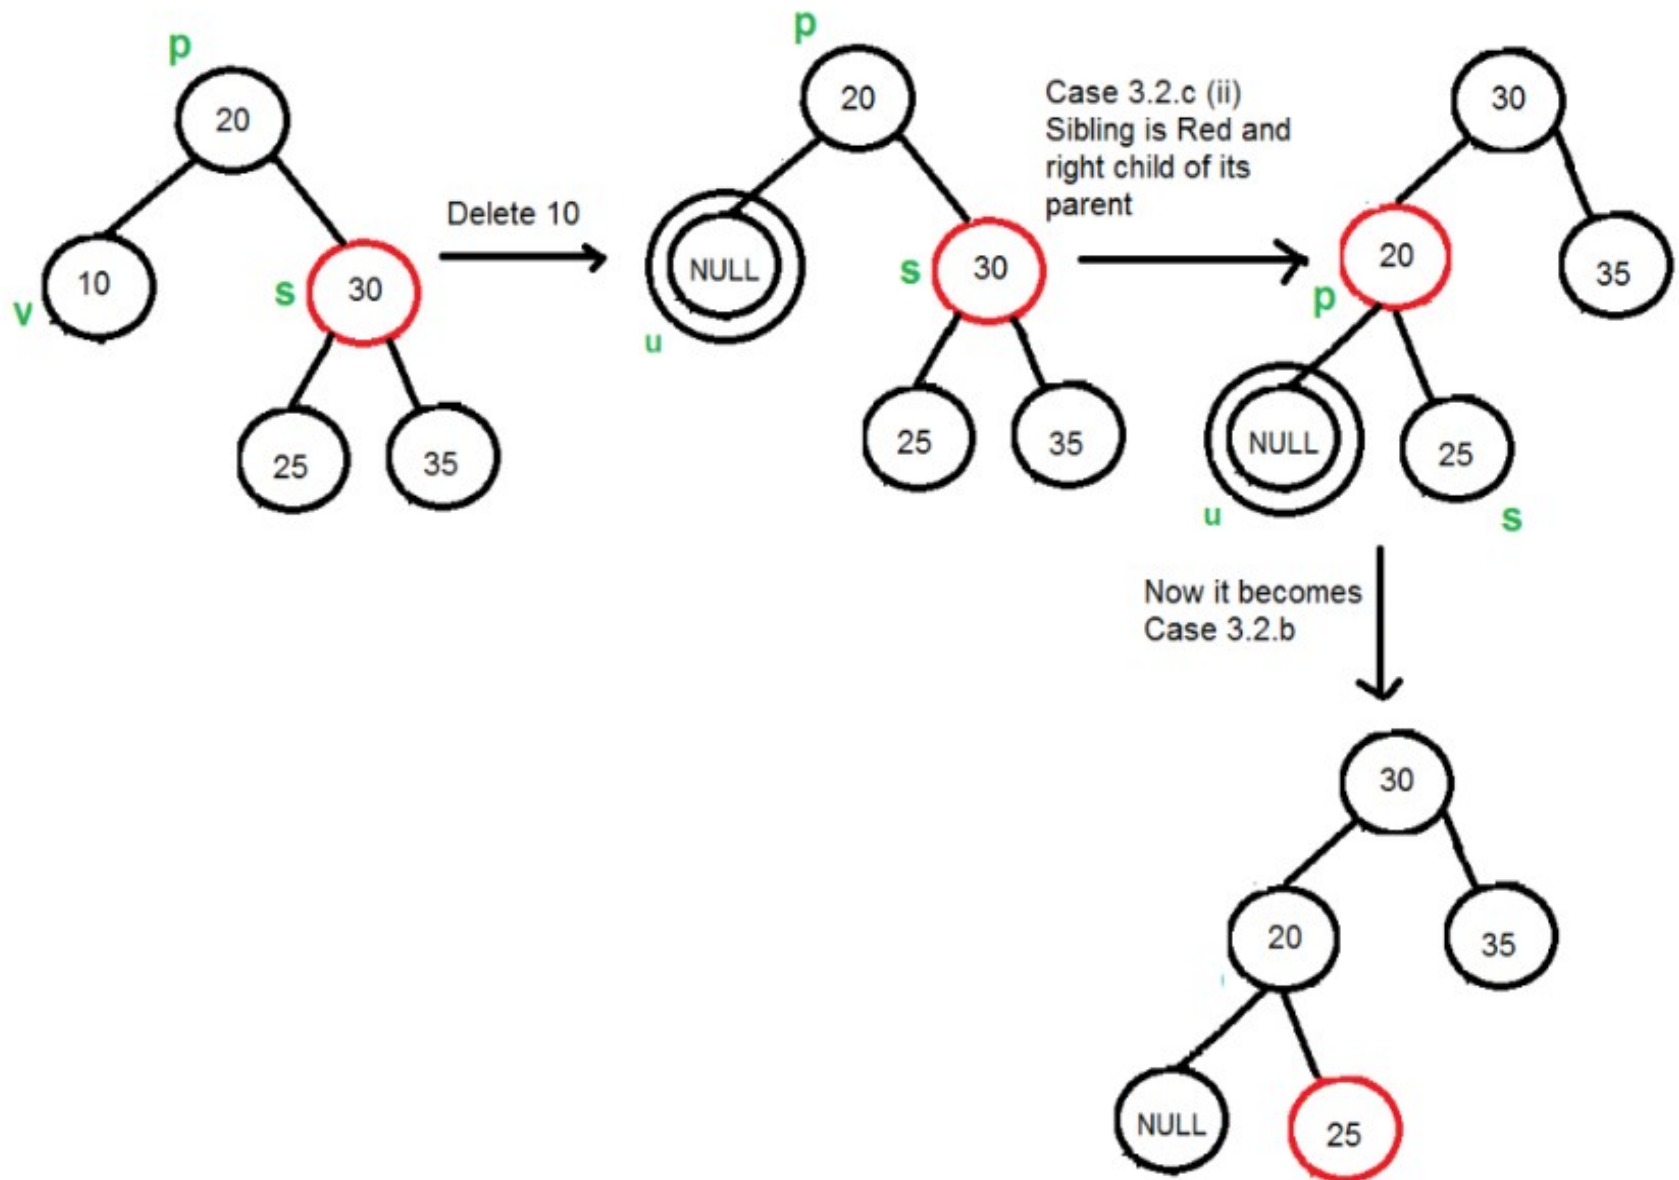




- **(b): If sibling is black and its both children are black**, perform recoloring, and recur for the parent if parent is black.
- In this case, if parent was red, then we didn't need to recur for parent, we can simply make it black (red + double black = single black)



- **(c): If sibling is red**, perform a rotation to move old sibling up, recolor the old sibling and parent. The new sibling is always black . This mainly converts the tree to black sibling case (by rotation) and leads to case (a) or (b).
- This case can be divided in two subcases.
  - (i)** Left Case (s is left child of its parent). This is mirror of right right case shown in below diagram. We right rotate the parent
  - (iii)** Right Case (s is right child of its parent). We left rotate the parent p.





- 
- **3.3)** If  $u$  is root, make it single black and return (Black height of complete tree reduces by 1).

# Red-black Tree Deletion

- First use the standard BST tree deletion algorithm
  - If the node to be deleted is replaced by its successor/predecessor (if it has two non-null children), consider the deleted node's data as being replaced by its successor/predecessor's, and its color remaining the same
    - The successor/predecessor node is then removed
- Let  $y$  be the node to be removed
- If the removed node was red, no property could get violated, so just remove it.
- Otherwise, remove it and call the tree-fix algorithm on  $y$ 's child  $x$  (the node which replaced the position of  $y$ )
  - Remember, the removed node can have at most one real (non-null) child
  - If it has one real child, call the tree-fix algorithm on it
  - If it has no real children (both children are null), Note that this child may be a (black) pretend (null) child


# Fixing a red-black Tree

---

- The tree-fix algorithm considers the parameter ( $x$ ) as having an “*extra*” *black token*
  - This corrects the violation of property 4 caused by removing a black node
- If  $x$  is red, just color it black
- But if  $x$  is black then it becomes “doubly black”
  - This is a violation of property 1
  - The extra black token is pushed up the tree until
    - a red node is reached, when it is made black
    - the root node is reached or
    - it can be removed by rotating and recoloring

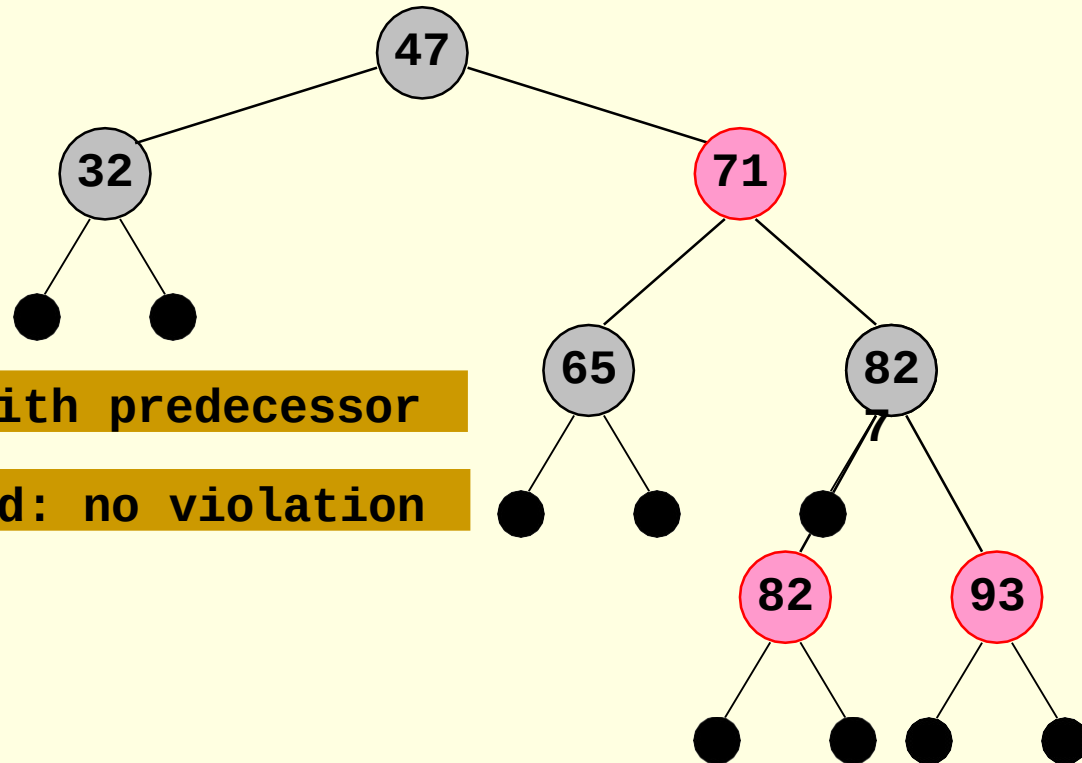
# Deletion Example 1

Delete 87



# Deletion Example 1

Delete 87




Replace data with predecessor

Predecessor red: no violation

# Deletion Example 2

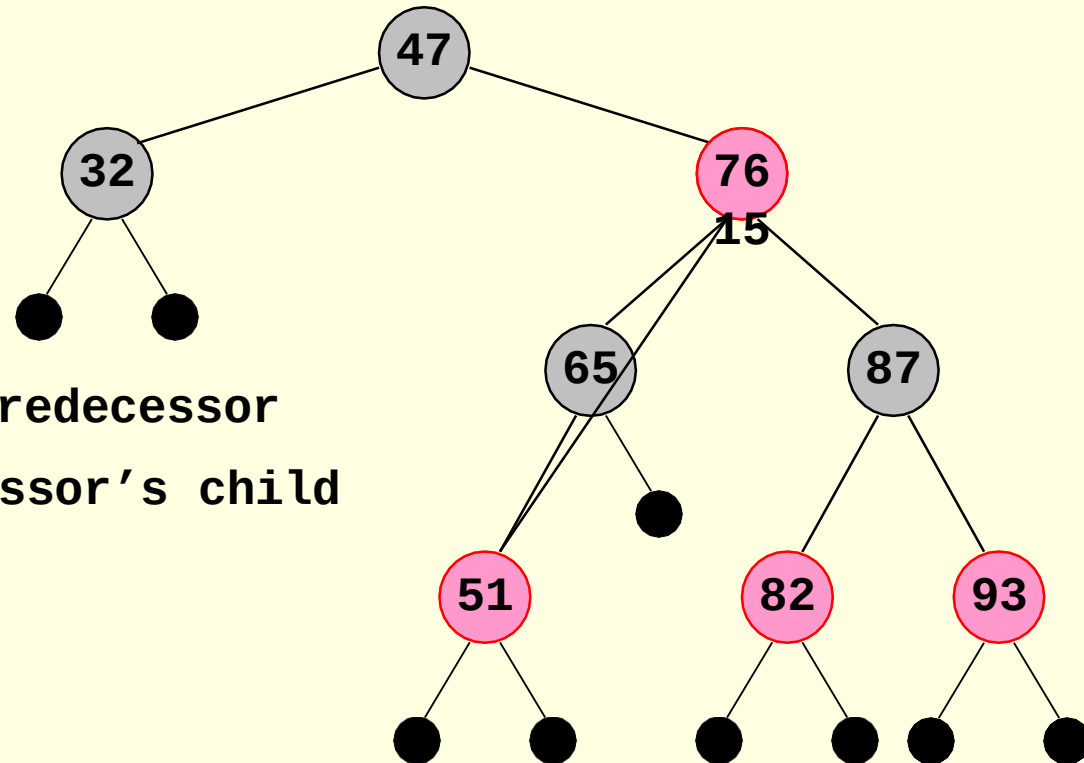
Delete 71



# Deletion Example 2

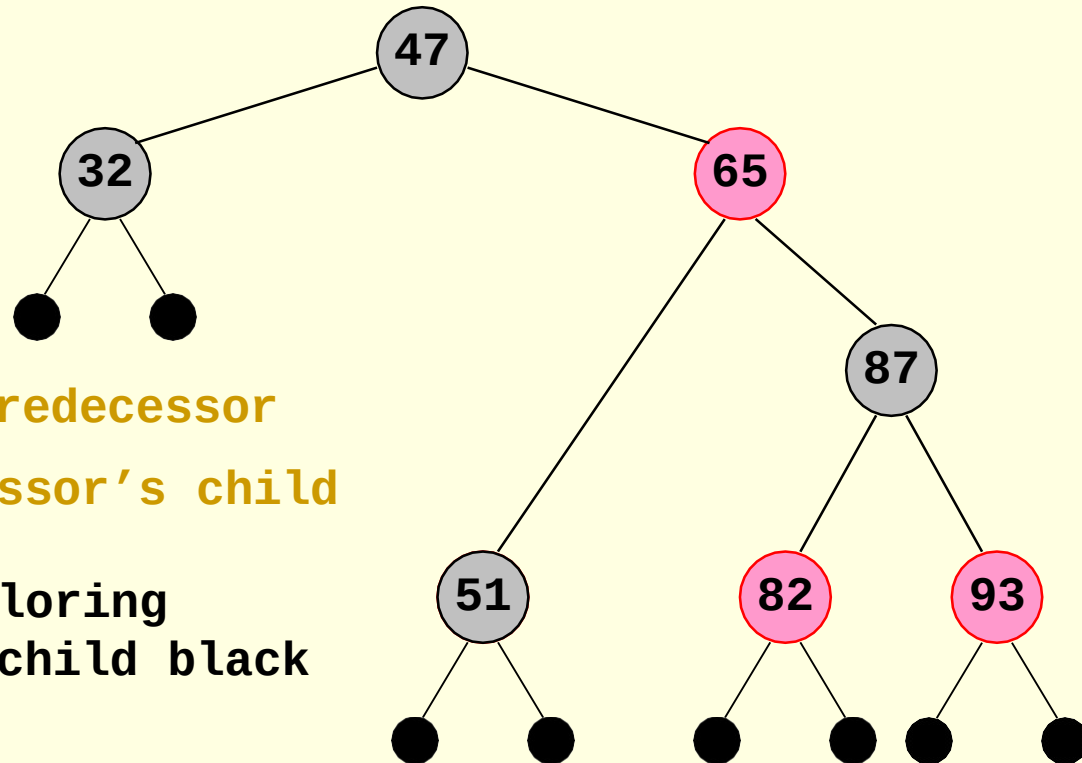
Delete 71

Replace with predecessor  
Attach predecessor's child



# Deletion Example 2

Delete 71



Replace with predecessor


Attach predecessor's child

Fix tree by coloring  
predecessor's child black




# Deletion Example 3

Delete 32



# Deletion Example 3

Delete 32



Identify x – the removed  
node's left child

Remove target node

Attach x to parent of  
target

# Deletion Example 3


Delete 32

Identify x - the removed  
node's left child

Remove target node

Attach x to parent of  
target

Call `rbTreeFix` on x



# RB Tree Deletion Algorithm

```
TreeNode<T> rbDelete(TreeNode<T> root,TreeNode<T> z)
```

```
//return new root, z contains item to be deleted
```

```
{
    TreeNode<T> x,y;
    // find node y, which is going to be removed
    if (z.getLeft() == null || z.getRight() == null)
        y = z;
    else {
        y = successor(z); // or predecessor
        z.setItem(y.getItem()); // move data from y to z
    }

    // find child x of y
    if (y.getRight() != null)
        x = y.getRight();
    else
        x = y.getLeft();

    // Note x might be null;
    // create a pretend node
    if (x == null) {
        x = new TreeNode<T>(null);
        x.setColor(black);
    }
}
```

# RB Tree Deletion Algorithm

```
x.setParent(y.getParent()); // detach x from y
if (y.getParent() == null)
    // if y was the root, x is a new root
    root = x;
else
    // Attach x to y's parent
    if (y == y.getParent().getLeft()) // left child
        y.getParent().setLeft(x);
    else
        y.getParent().setRight(x);

if (y.getColor() == black)
    root=rbTreeFix(root,x);

if (x.getItem() == null) // x is a pretend node
    if (x==x.getParent().getLeft())
        x.getParent().setLeft(null);
    else
        x.getParent().setRight(null);


return root;
}
```

# Deletion Example 3

## (continued)


After deleting 32,  
x is a node with  
black token

Identify y, x's sibling  
Make y black and y's parent  
red  
Left rotate x's parent



# Deletion Example 3

After deleting 32,  
x is a node with  
black token



Identify y, x's sibling


Make y black and  
y's parent red

Left rotate x's parent

Identify y - x's new  
sibling

# Deletion Example 3

After deleting 32,  
x is a node with  
black token



Identify y, x's sibling

Make y black and  
y's parent red

Left rotate x's parent

Identify y - x's new sibling

Color y red

Assign x it's parent, and color it black




## Tree Fix algorithm cases: case (1) x is

red

---

- The simplest case
- x has a black token and is colored red, so just color it black and remove token and we are done!
- In the remaining cases, assume x is black (and has the black token, i.e., it's double black)

# Tree Fix algorithm cases: case (2) ~~x's sibling is red~~



Colors of y and z were swapped


Colors of x and y were swapped

## Remarks:

- the roots of subtrees C and D are black
- the second is the symmetric case, when x is the right child
- in the next step (case (3) or (4)) the algorithm will finish!

# Tree Fix algorithm cases: case (3)

x's sibling is black and both nephews are black




Remarks:

- nephews are roots of subtrees C and D

# Tree Fix algorithm cases: case (4)

**x's sibling is black and at least one nephew is red**



Colors of y and z were swapped. Far nephew is colored black and black token is removed.

Colors of z and y were swapped. Far nephew is colored black and black token is removed.

Remarks:

- in this case, the black token is removed completely
- if the “far” nephew is black (subcase (i)), rotate its parent, so that

# RB Trees efficiency

- All operations work in time  $O(\text{height})$
- and we have proved that height is  $O(\log n)$
- hence, all operations work in time  $O(\log n)$ ! – much more efficient than linked list or arrays implementation of sorted list!

<b>Sorted List</b>	<b>Search</b>	<b>Insertion</b>	<b>Deletion</b>
with arrays	$O(\log n)$	$O(n)$	$O(n)$
with linked list	$O(n)$	$O(n)$	$O(n)$
with RB trees	$O(\log n)$	$O(\log n)$	$O(\log n)$